

Just to note, this is work in progress (unfinished)

Summary of Asmcodes series

Introduction

Between November 2015 and December 2016, a number of cryptographic algorithms specifically designed for software implementation were selected for optimization in C and x86 assembly code. The algorithms consisted of block and stream ciphers, cryptographic hash functions and some modular arithmetic functions necessary for asymmetric key exchange and verification of digital signatures.

The purpose of this work was to evaluate the suitability of such algorithms for resource constrained environments which could be anything from a small microcomputer to a block of memory.

Although the vast majority of low resource devices do not use the x86 architecture, I think the results presented here could offer valuable insight for anyone searching for cryptographic algorithms suitable for resource constrained environments . Most of the assembly codes were a joint effort with Peter Ferrie. [1]

Desirable properties

The criterion for selection was very simple:

- Patent free
- Moderate to high level of security
- Small space required in for both ROM/RAM with RAM being a lesser consideration.
- Simple implementation

Many ciphers despite their popularity don't meet the criteria. Some ciphers are strong but complicated to implement or have large resource requirements.

Blowfish for example while considered immune to many cryptographic attacks uses large lookup tables to create key dependent s-boxes which cannot be calculated at runtime. This requires both large amounts of RAM and ROM therefore making it unsuitable for low resource devices. The same is true for ciphers such as MARS, Camellia, CAST and cryptographic hash algorithms like Streebog or Whirlpool.

The stream cipher HC-256 although incredibly small in ROM takes as much RAM as Blowfish since it uses a similar function for generating key stream. The assembly implementation for example uses 20KB of RAM for key generation and 8KB lookup tables.

For the windows operating system, allocating 20KB requires a small amount of code to perform a *page probe* in order to prevent exceptions occurring.

MARS also uses large lookup tables and again can't be calculated at runtime. In the following table, you will see a full list of block and stream ciphers that were examined for implementation. The algorithm rows highlighted in grey were implemented but due to patents will not be made available in final library.

Algorithm	Type	Implemented	Low Resource
AES-256	Block	Yes	Yes
Serpent-256	Block	Yes	Yes
Twofish-256	Block	Yes	Yes
CAST-256	Block	No	
ThreeFish-256	Block	Yes	Yes
Noekeon	Block	Yes	Yes
Speck	Block	Yes	Yes
Camellia-256	Block	No	
MARS	Block	No	
Kalyna	Block	No	
Bel-T-256	Block	Yes	Yes
Kuznyechik	Block	Yes	Yes
Blowfish	Block	Yes	No
DES	Block	Yes	No
RC5	Block	Yes	Yes
RC6	Block	Yes	Yes
ChaCha20	Stream	Yes	Yes
HC-256	Stream	Yes	Yes/No
Salsa20	Stream	Yes	Yes
Rabbit	Stream	Yes	Yes
SHA3	Hash	Yes	Yes
BLAKE2	Hash	Yes	Yes
SHA2	Hash	Yes	Yes
Half Sip Hash	Hash	Yes	Yes
Chaskey	MAC	Yes	Yes
Poly1305	MAC	Yes	Yes
CubeHash	Hash	Yes	Yes
MD4	Hash	Yes	Yes
MD5	Hash	Yes	No
SHA1	Hash	Yes	Yes

Based on the results, AES-256 is favourable as block cipher with Salsa20 as stream cipher and SHA3 as hash. The security of AES-256 is constantly being examined since it was made a standard in 2000 by NIST. Salsa20 was examined extensively as part of the eSTREAM competition and although similar to ChaCha20 by the same author, the latter has probably not been scrutinized as much as Salsa (open to correction).

Symmetric Encryption

The table below shows the differences between ciphers in byte size for both compiler generated and hand written x86 assembly. X64 assembly would undoubtedly result in larger code but are not implemented. Percentages shown here are approximations of hand written assembly versus assembly generated by compiler.

Cipher	Type of cipher	ROM (x86 ASM) Bytes	ROM (MSVC) Bytes
DES	Block	1038	1381
TWOFISH-256	Block	610	1242
SERPENT-256	Block	530	1027
KUZNYECHIK	Block	615	928
AES-256	Block	377	827
THREEFISH-256	Block	371	747
RABBIT	Stream	457	700
BEL-T	Block	490	656
NOEKEON	Block	292	431
SALSA20	Stream	245	436
CHACHA20	Stream	241	424
RC6	Block	247	408
HC-256	Stream	272	*354
RC5	Block	167	237
SPECK	Block	105	139

*= requires additional hidden code to increase stack size

Cryptographic Hash, MAC and checksum algorithms

The main algorithms selected here were from the SHA-3 competition although HMAC-SHA-2 was also included. MD4, MD5 and SHA-1 were merely included for comparison but were not seriously considered for use in the end.

Algorithm	Type	ROM (x86 ASM) Bytes	ROM (MSVC) Bytes
SHA2-256	Hash	653	1324
MD5	Hash	685	1202
BLAKE2-256	Hash	509	1027
SHA3	Hash	460	879
MD4	Hash	396	535
SHA-1	Hash	447	500
Poly1305	MAC	332	507
Chaskey	MAC	234	346
Cube Hash	Hash	201	448
Half Sip Hash 32	PRF	142	261
CRC32	Checksum	50	
INTEL CRC32	Checksum	46	

Based on implementations, SHA-3 has advantage for being the standard and will undoubtedly receive more scrutiny than other algorithms listed. Cube Hash is incredibly compact but also slow unless using a vectorized version. The nice thing about Cube Hash is the ability to tweak parameters and would be useful for very small applications.

Key Exchange

Due to the simplicity of factoring based public key exchange, only RSA and Diffie-Hellman-Merkle were examined. Elliptic Curve and Lattice based methods result in much more code although could be considered as an alternative if space isn't important.

The main function used for both methods of key exchange is Modular Exponentiation.

Obviously using RSA is much faster but I can't tell you which is more secure.

Final Results

Two configurations were selected.

	Algorithm	Size (MSVC)	Size (asm)
Symmetric	AES-256	760	377
Key Exchange	RSA-2048	844	138
MAC	SHA3	879	460
Total		2483	975

	Algorithm	Size (MSVC)	Size (asm)
Symmetric	Salsa20	436	245
Key Exchange	RSA-2048	867	138
MAC	Cube Hash	448	201
Total		1751	584

	Algorithm	Size (MSVC)	Size (asm)
Symmetric	ChaCha20	424	241
Key Exchange	RSA-2048	844	138
MAC	Poly1305	507	339
Total		1775	718

Future Work

The CAESAR competition is expected to be finalized by December 2017 at which time this document will hopefully be updated to reflect a number of algorithms chosen.

Summary

The numbers presented in all tables are approximations since further modifications would decrease the sizes significantly. For example, if using CTR mode for symmetric encryption, the decryption functionality can be removed for some algorithms which include inverse tables.

The parameters of MAC can be fixed. Multiple calls can be reduced to one. Multiple functions can be “glued” together. Many versions could shave hundreds of bytes from at least the assembly code level.

Although it's entirely possible to write offset independent code using pure C and such code is easily ported to other architectures, the compiler could never possibly compete with experienced assembly programmers given the current compilers available.

Depending on the design of the algorithms, a compiler does come close to hand-written assembly but is usually not capable of reducing as much as assembly programmer.

References

1. Homepage of Peter Ferrie
<http://pferrie.host22.com/>
2. AES: the Advanced Encryption Standard
<https://competitions.cr.yp.to/aes.html>
3. eSTREAM: the ECRYPT Stream Cipher Project
<https://competitions.cr.yp.to/estream.html>
4. SHA-3: a Secure Hash Algorithm
<https://competitions.cr.yp.to/sha3.html>
5. CAESAR: Competition for Authenticated Encryption: Security, Applicability, and Robustness
<https://competitions.cr.yp.to/caesar.html>